

Lab 1

Deadline: January 27, 23:59.

Welcome to the first lab! Starting off, please familiarize yourself with the Github repository (<https://github.com/tollefj/TDT4310-spring-2023>). You will need to clone it to be able to work on the labs.

All exercises/labs will follow this format:

- A markdown file (also supplied as a PDF) – which you are currently reading – with questions and tasks relevant to the current lab.
- Server code that you will implement to improve the keyboard, found under `labs/your_implementations/lab_n.py`.

It is highly recommended to complete the initial questions before moving on to the implementation. They should not be time consuming. Please use the discussion boards

You will deliver the following

A zipped file with your NTNU user name, containing two files:

- A document (any sensible format, but PDF preferred) which answers the tasks in this document
- A copy of your `lab_1.py` file at `labs/your_implementations/lab_1.py`

1) Tokenization

Consider the following sentence:

My first car (1974 Ford Pinto) was a trash-can on wheels...

1. How would *you* tokenize this sentence into words? No coding required (nor is there a “correct” answer)
2. If you were to type the sentence on your phone, what would you expect the next prediction to be after typing “My first car”?
3. What are some of your initial thoughts on the difficulty of next-word prediction in general?

2) Introduction to language modeling

Language modeling is, in short, the task of predicting the next word in a sentence. You may have heard of BERT, GPT, and other language models. A simpler language model can be implemented using **n-grams**. You will be implementing a *bigram* and *trigram* model in this lab. Before you tackle this task, some basic knowledge might come in handy.

With the sentence above, use NLTK to get:

1. Bigrams

2. Trigrams

We can continue the n-gram model infinitely (4-grams, 5-grams, ...)

3. What issues may occur if we select a large n value for a small corpus?
What would you guess to be ideal for a smart keyboard?

3) Introduction to word representations

Vectorization, word representations, or word embeddings, are a way to represent words or documents as vectors. Word vectors will be described in greater detail later in the course.

For now, we'll consider a very simple implementation based on the bag-of-words representation of a document.

With the sentence:

That that is is that that is not. Is that it? It is.

- NLTK allows us to create the following frequency distribution using `FreqDist`:
 - `{'is': 6, 'that': 5, 'not': 2, 'it': 2}`
- By considering the words as their indices in the vocabulary...
 - *that* $\rightarrow 0$, *is* $\rightarrow 1$, *not* $\rightarrow 2$, *it* $\rightarrow 3$
- The sentence can be represented as the values of each index:
 - `[5, 6, 2, 2]`

Answer the following:

1. What are your thoughts on the usefulness of this representation? Can you think of a way to improve it?
2. How could you use this technique to compare sentences with each other?

Implementation notes

As a final part of the lab, I want you to briefly discuss your approach to the implementation task. This could be things you had to learn, difficult parts of the lab, etc.